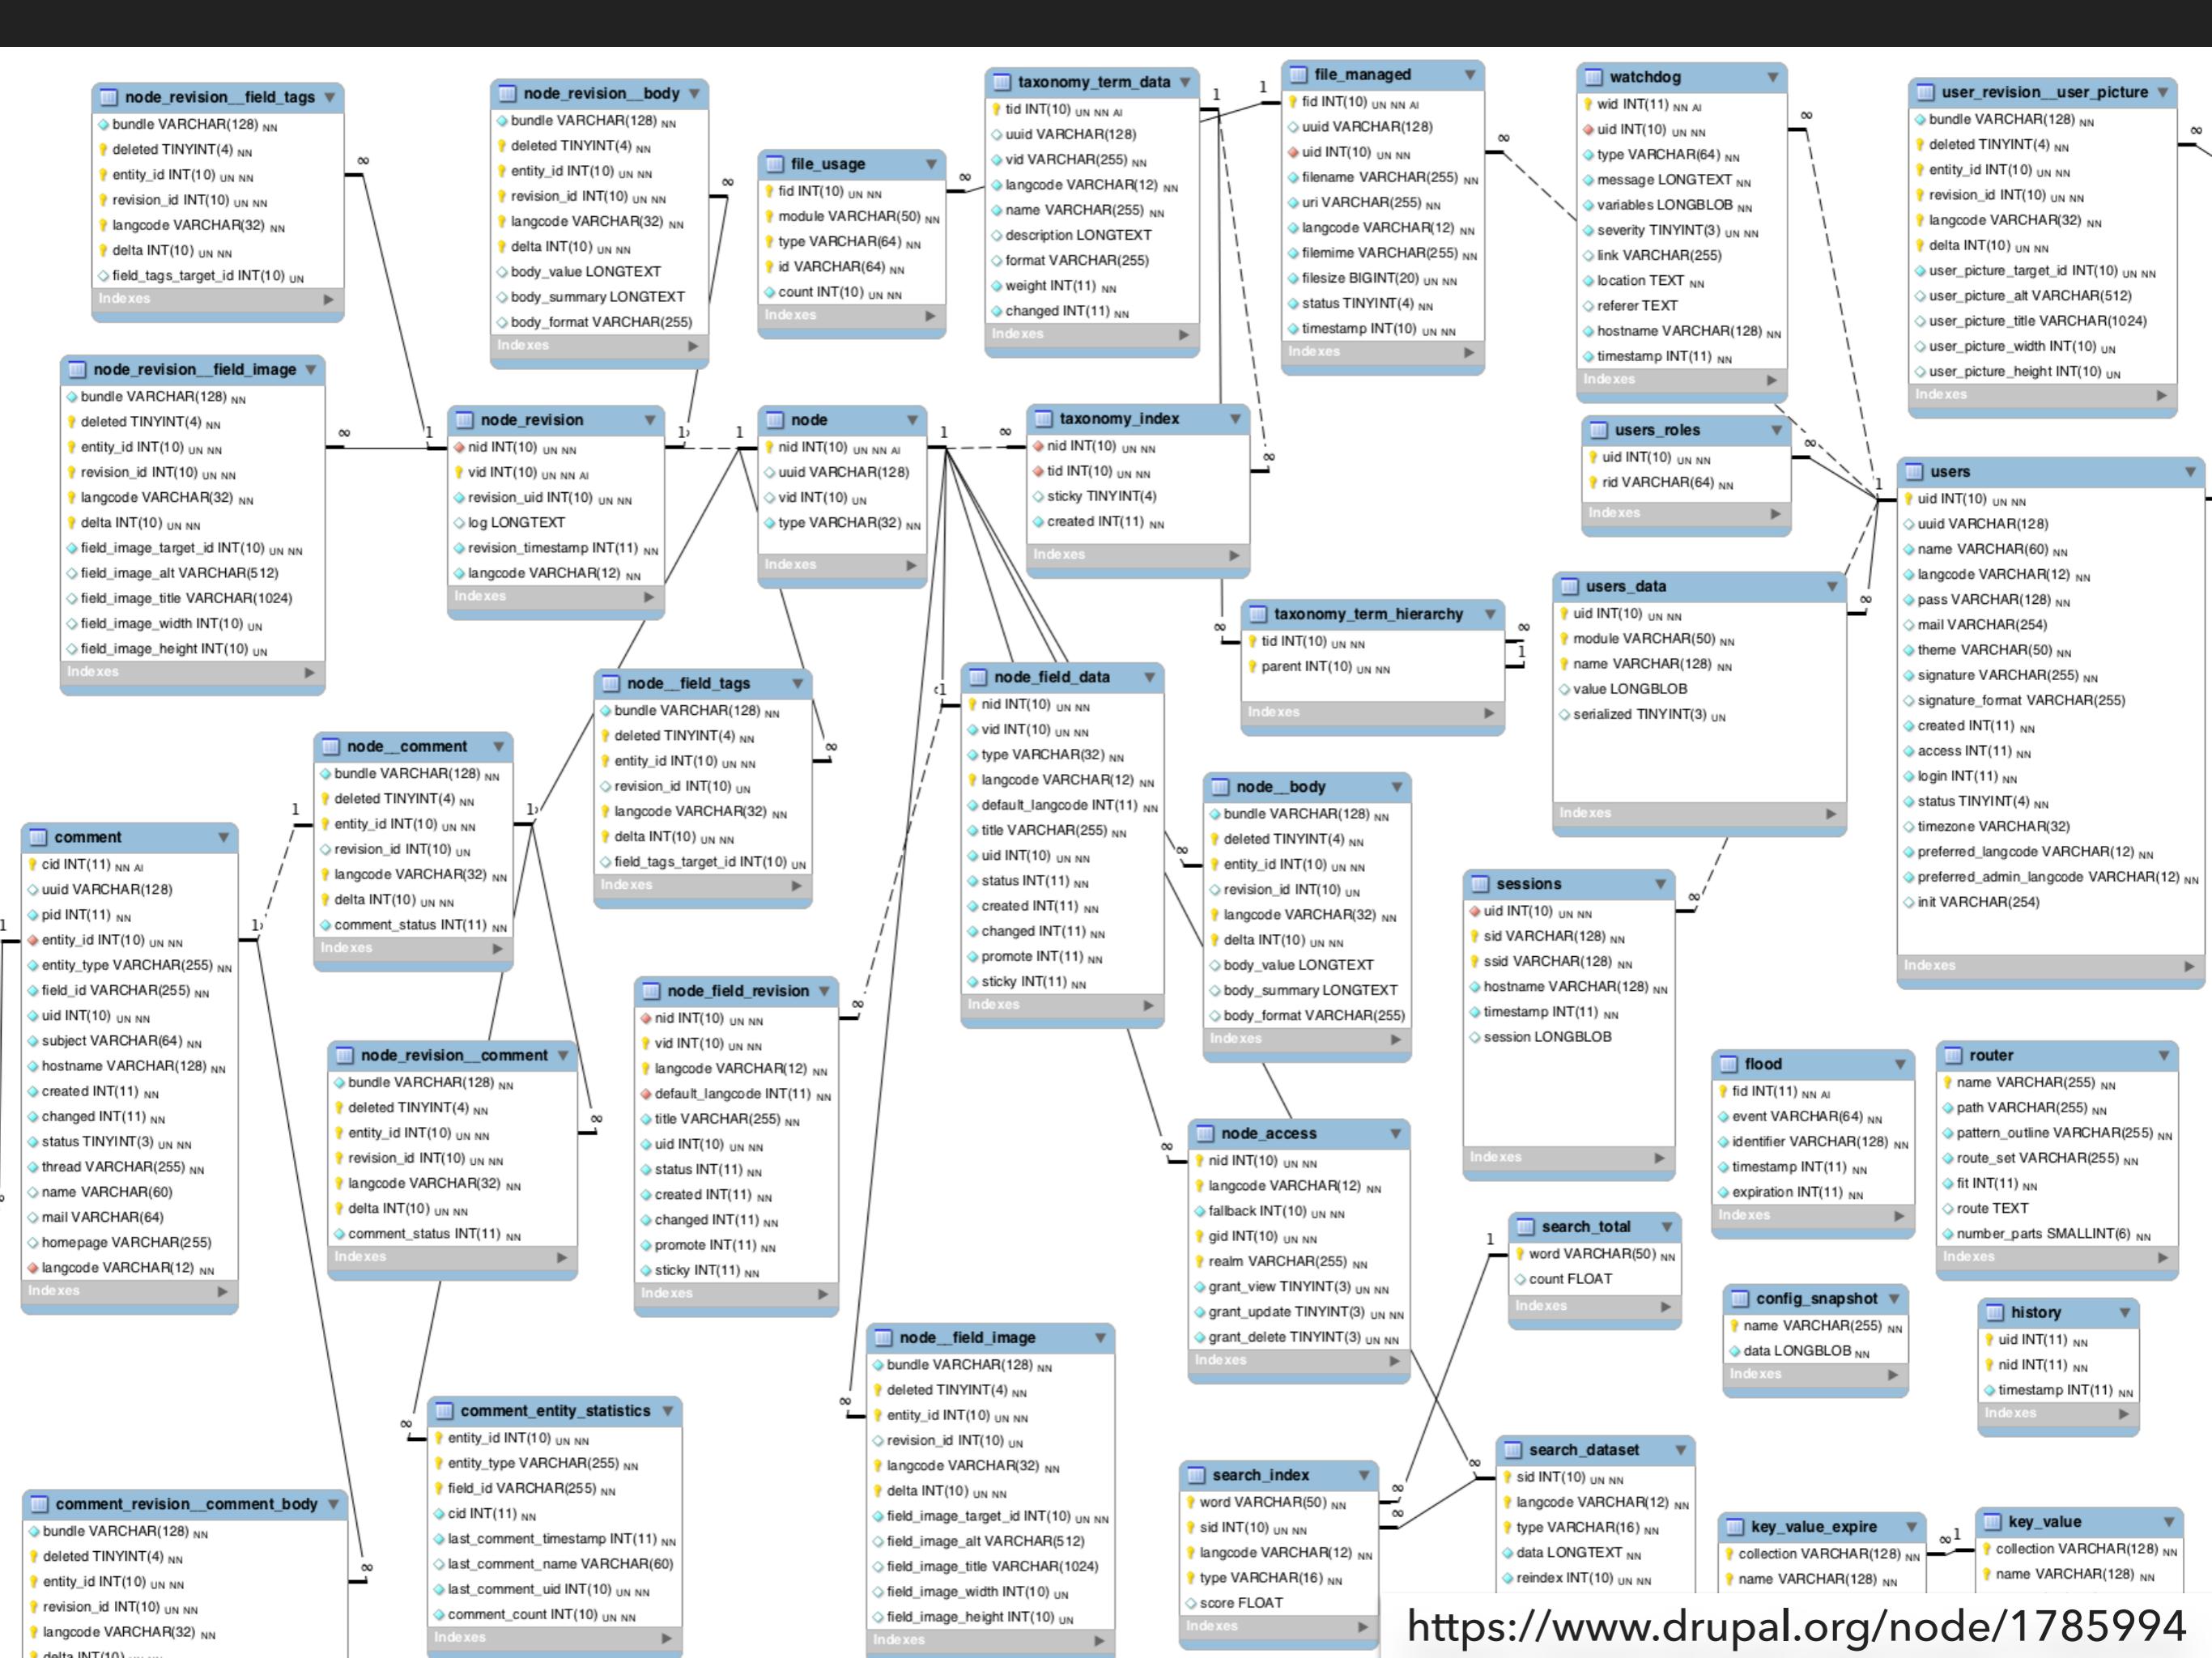


Question

If we update a Repair PO,
what else do we need to
save as part of that
transaction?

Transactional
boundaries are often
ad-hoc and create
sprawling dependencies

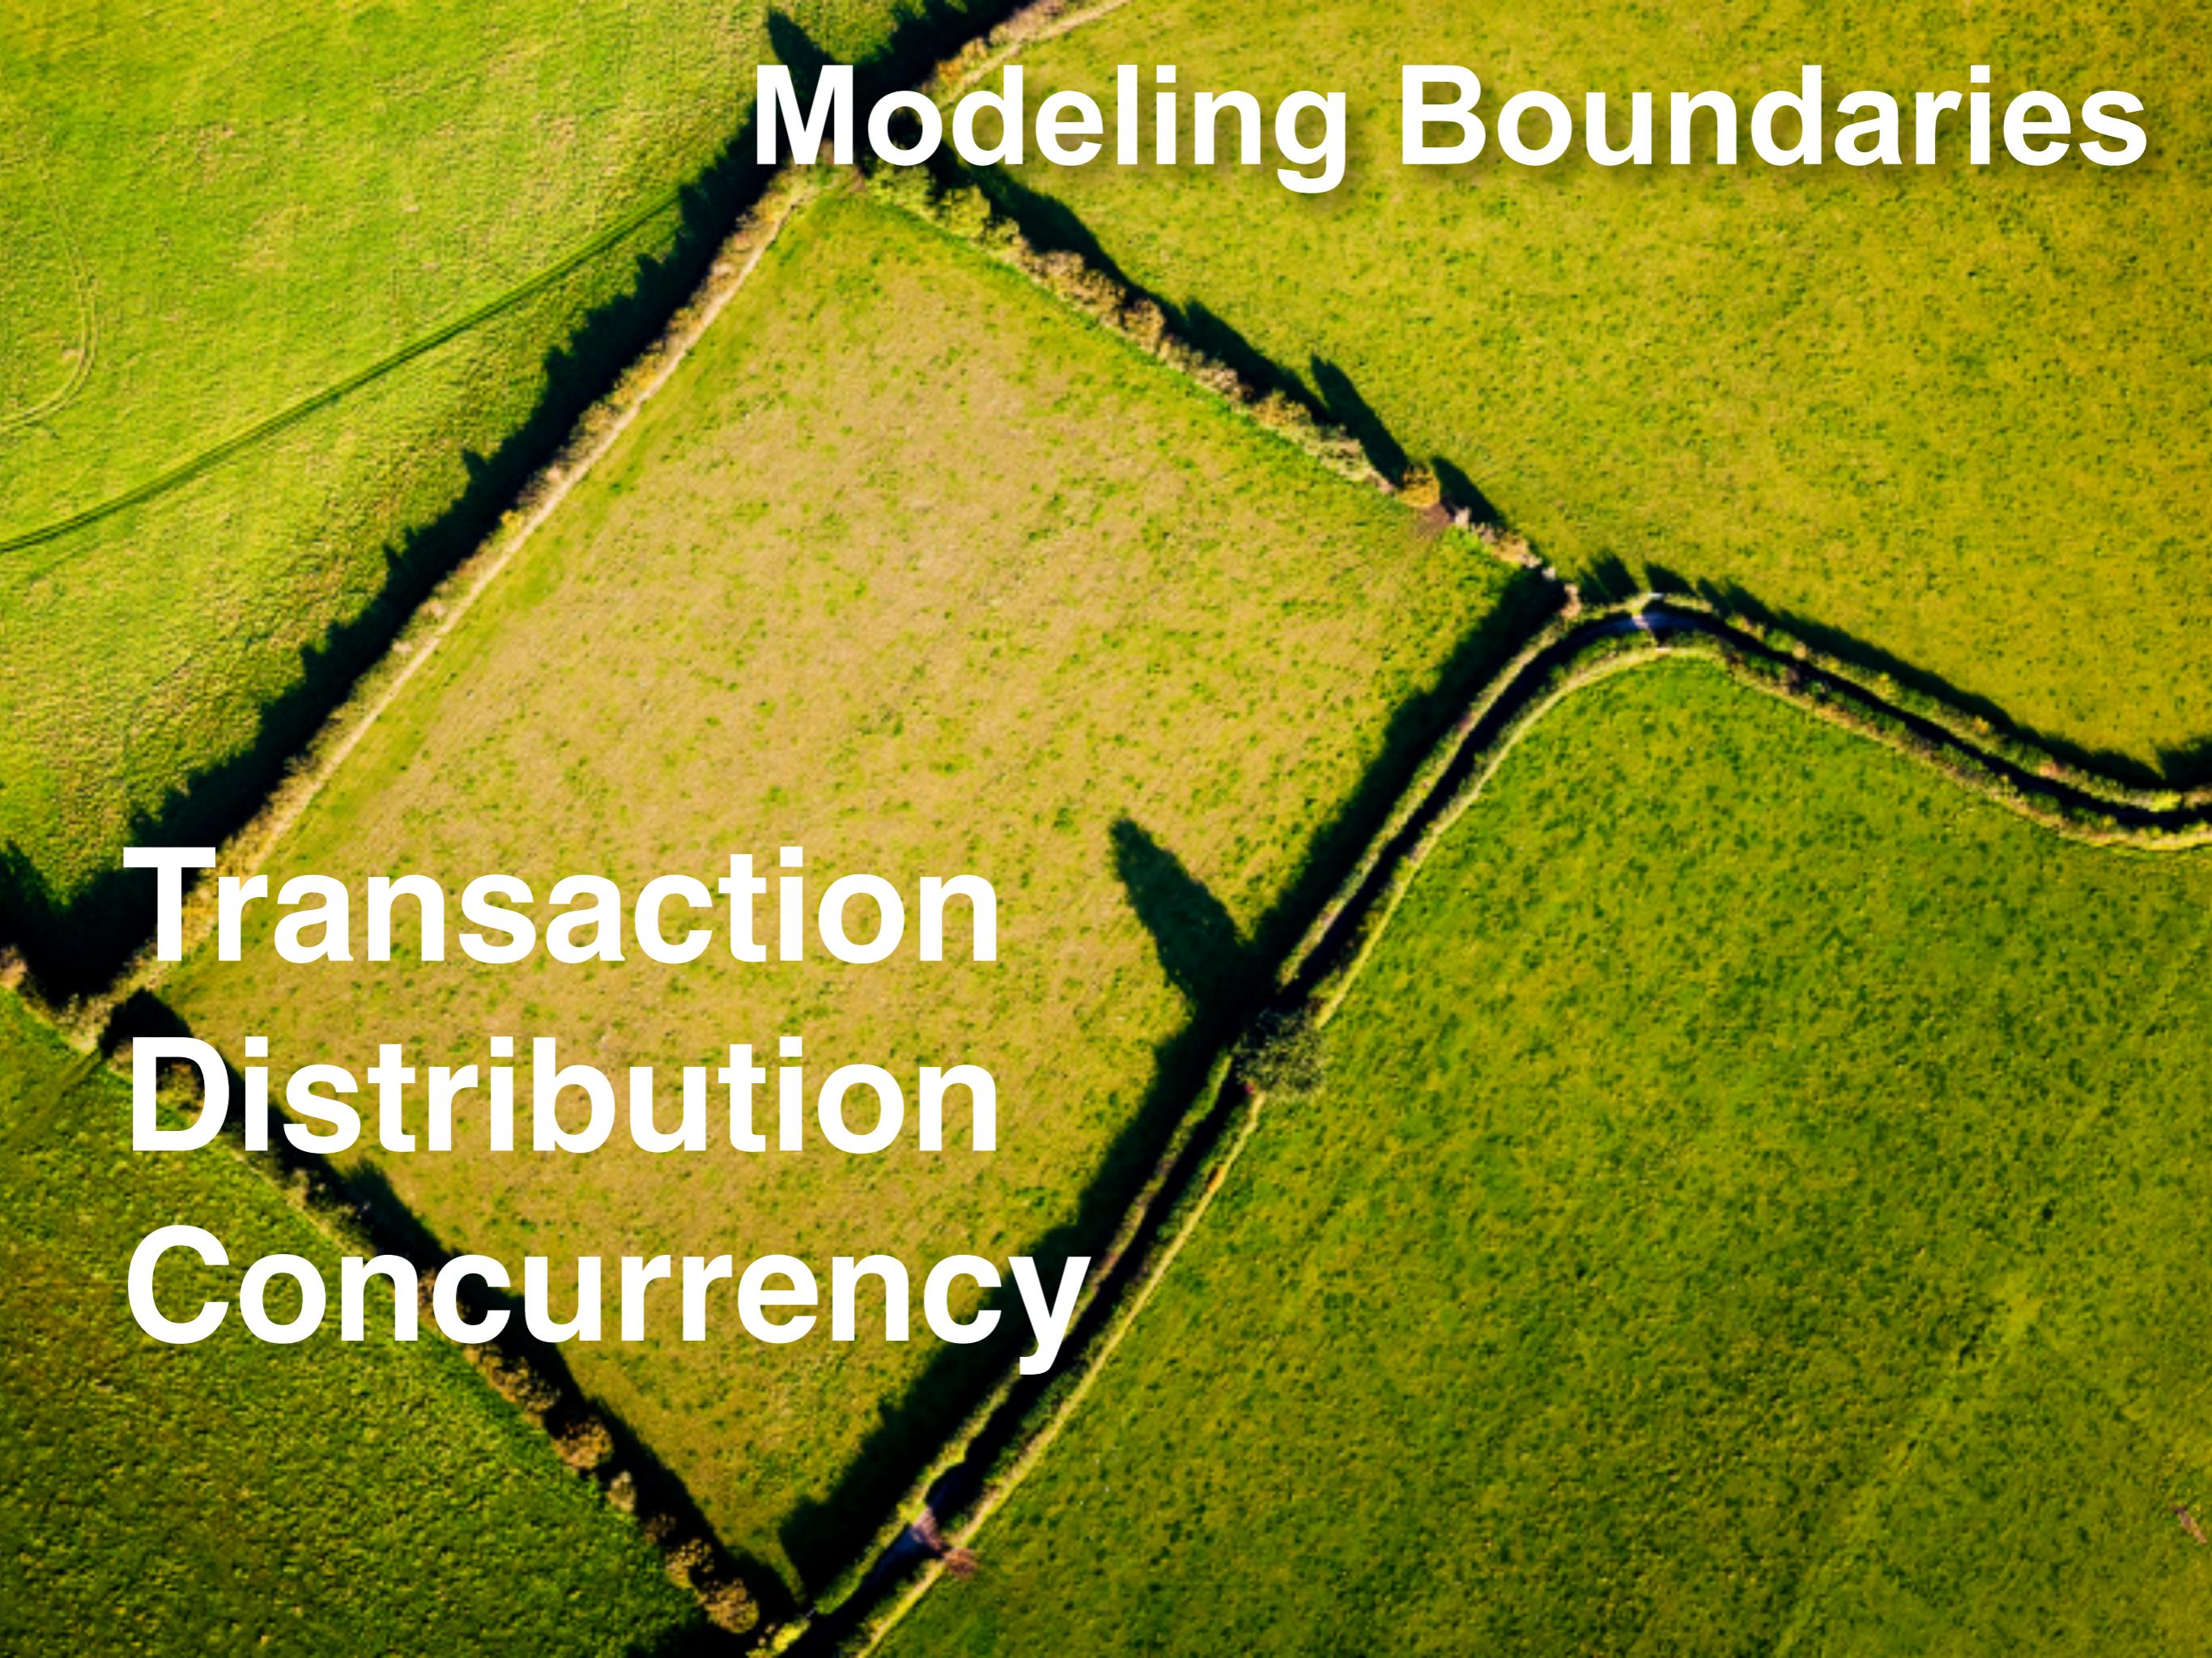


Building performant,
scalable, robust,
distributed systems
requires embracing
decoupling and latency
boundaries

Concurrency is tricky
to get right, and can
really negatively impact
performance



Aggregates

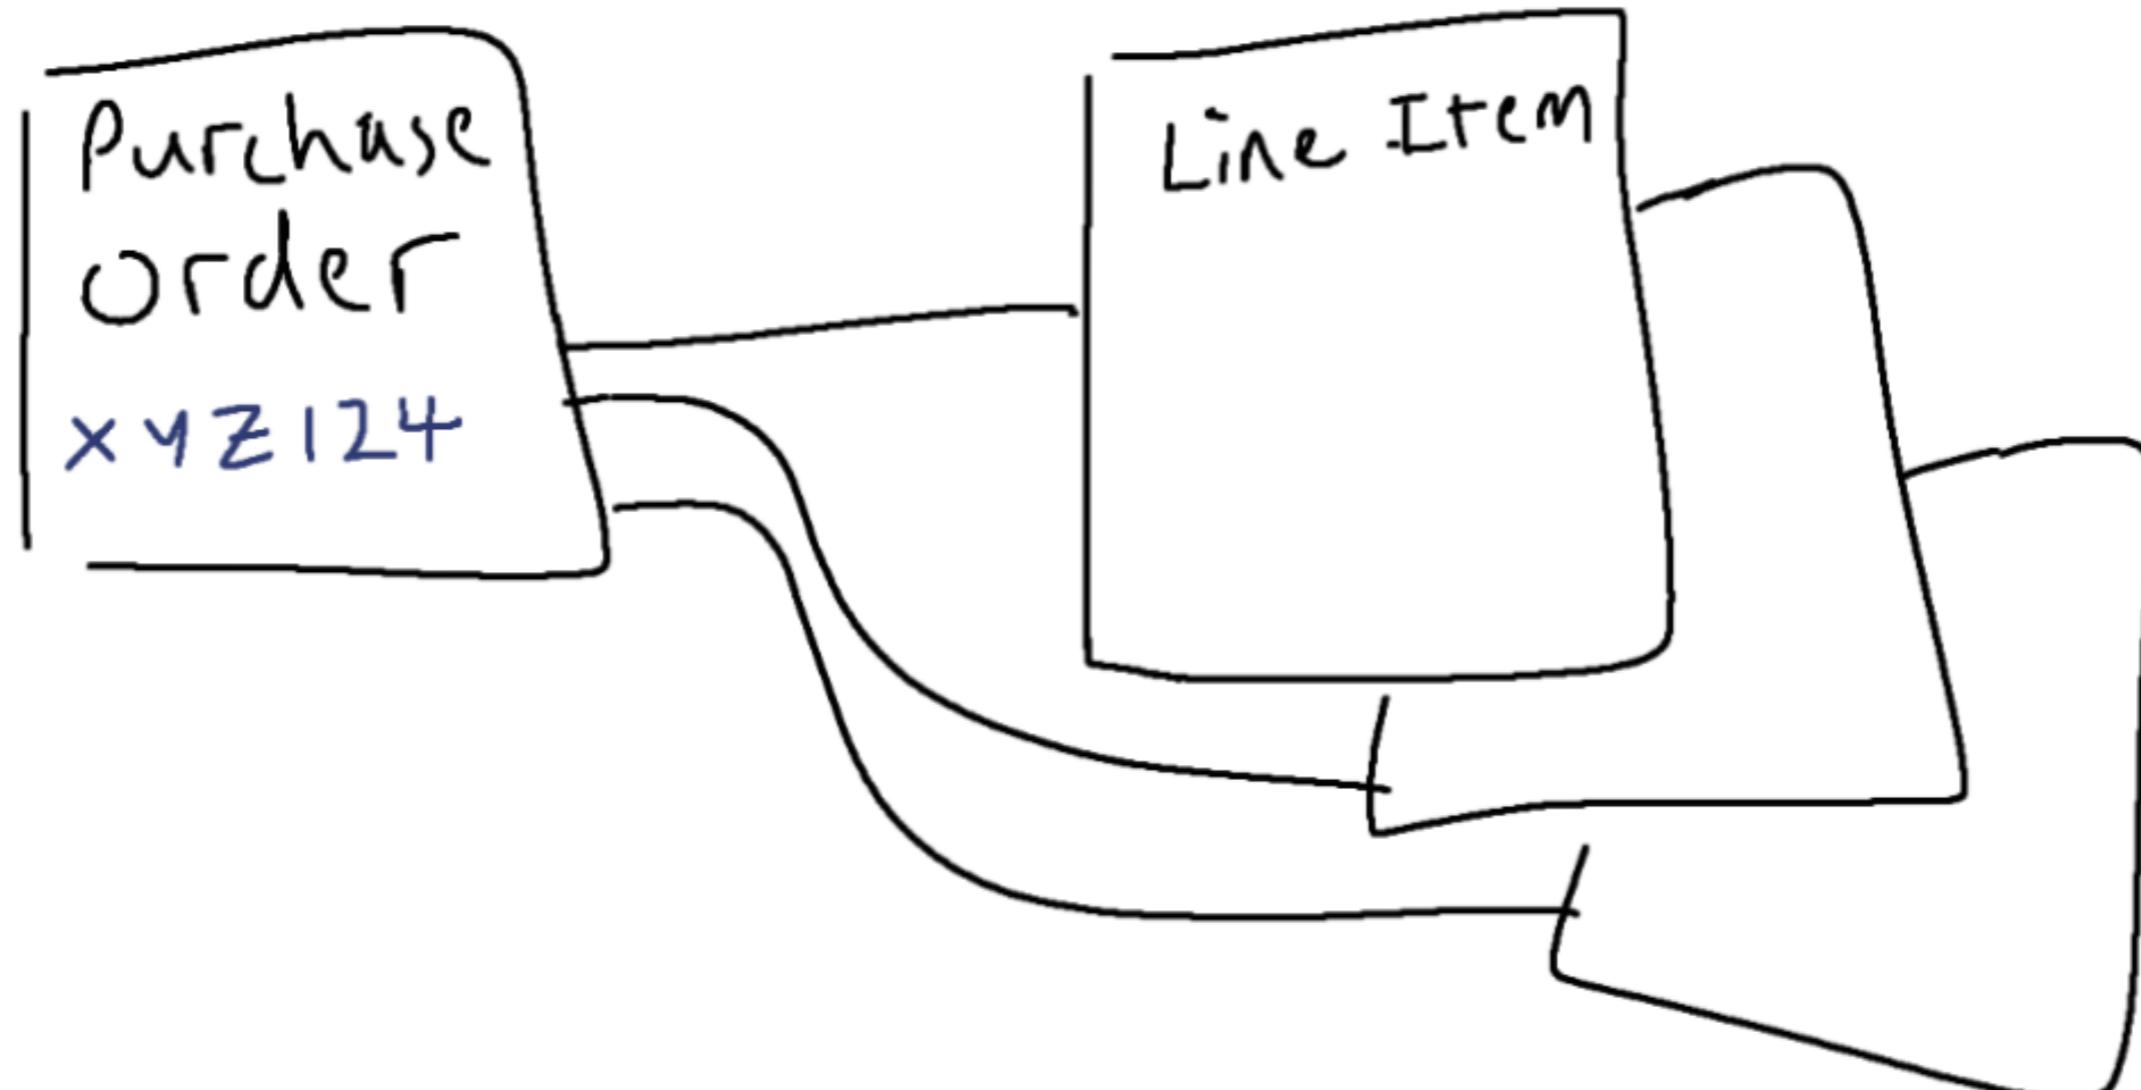
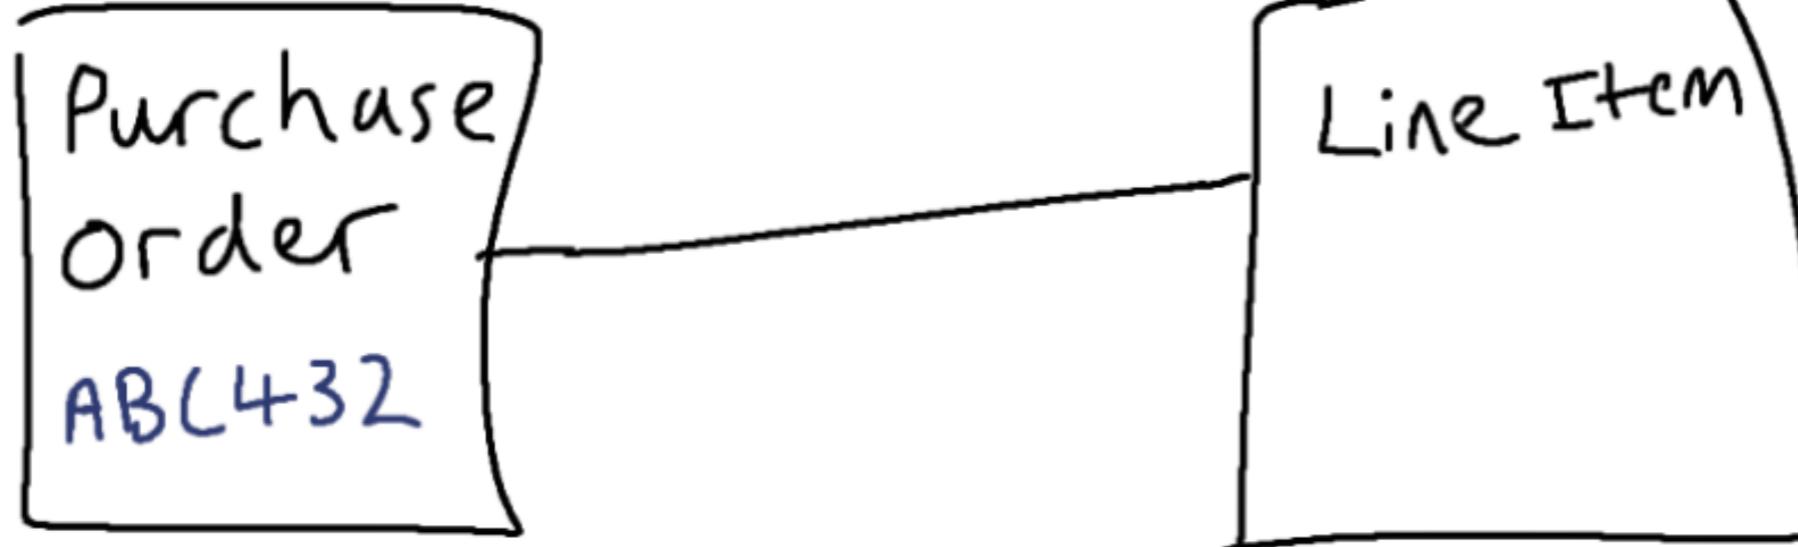
An aerial photograph of a rural landscape featuring several green fields. The fields are separated by various boundaries, including dirt paths, stone walls, and hedgerows. The lighting suggests it's either early morning or late afternoon, casting long shadows of the boundaries across the fields.

Modeling Boundaries

Transaction
Distribution
Concurrency

Aggregate

A cluster of domain objects
that can be treated as a
single unit





5 Guidelines for Aggregate Design

Guideline #1

Design small,
cohesive
aggregates

Avoid large, ad-hoc, sprawling
transactional, distribution and
concurrency boundaries

Guideline #2

Identify true invariants to
clarify aggregate
boundaries

Sample Invariants

Contract must always have valid:

- ClientCompany
- Product
- TermsAndConditions

Repair PO must always have at least one LineItem

What Needs to Change...

To separate RepairPurchaseOrders, Claims and Contracts into separate aggregates?

Guideline #3

Reference other
aggregates by
identity

Aggregate Root

An entity that serves as the identity holder, namesake, and gatekeeper for the aggregate.

Use the delete rule of thumb to assist you in identifying the root of an aggregate.

Use the Factory pattern to construct more complex aggregates to ensure invariants are enforced correctly.

Guideline #4

Only hold references to
aggregate roots
(i.e. limit interconnectedness)

New Invariants

- *A pending Contract cannot have any Claims.*
 - *Only one open Claim at a time per active Contract.*
 - *Sum of costs of all Claims for a Contract must never exceed the LimitOfLiability.*
-
- What needs to change in your design to make Claims part of the Contract aggregates again?
 - How well does your design handle relating each RepairPO to its associated Claim, especially when there are multiple Claims for a Contract?

Guideline #5

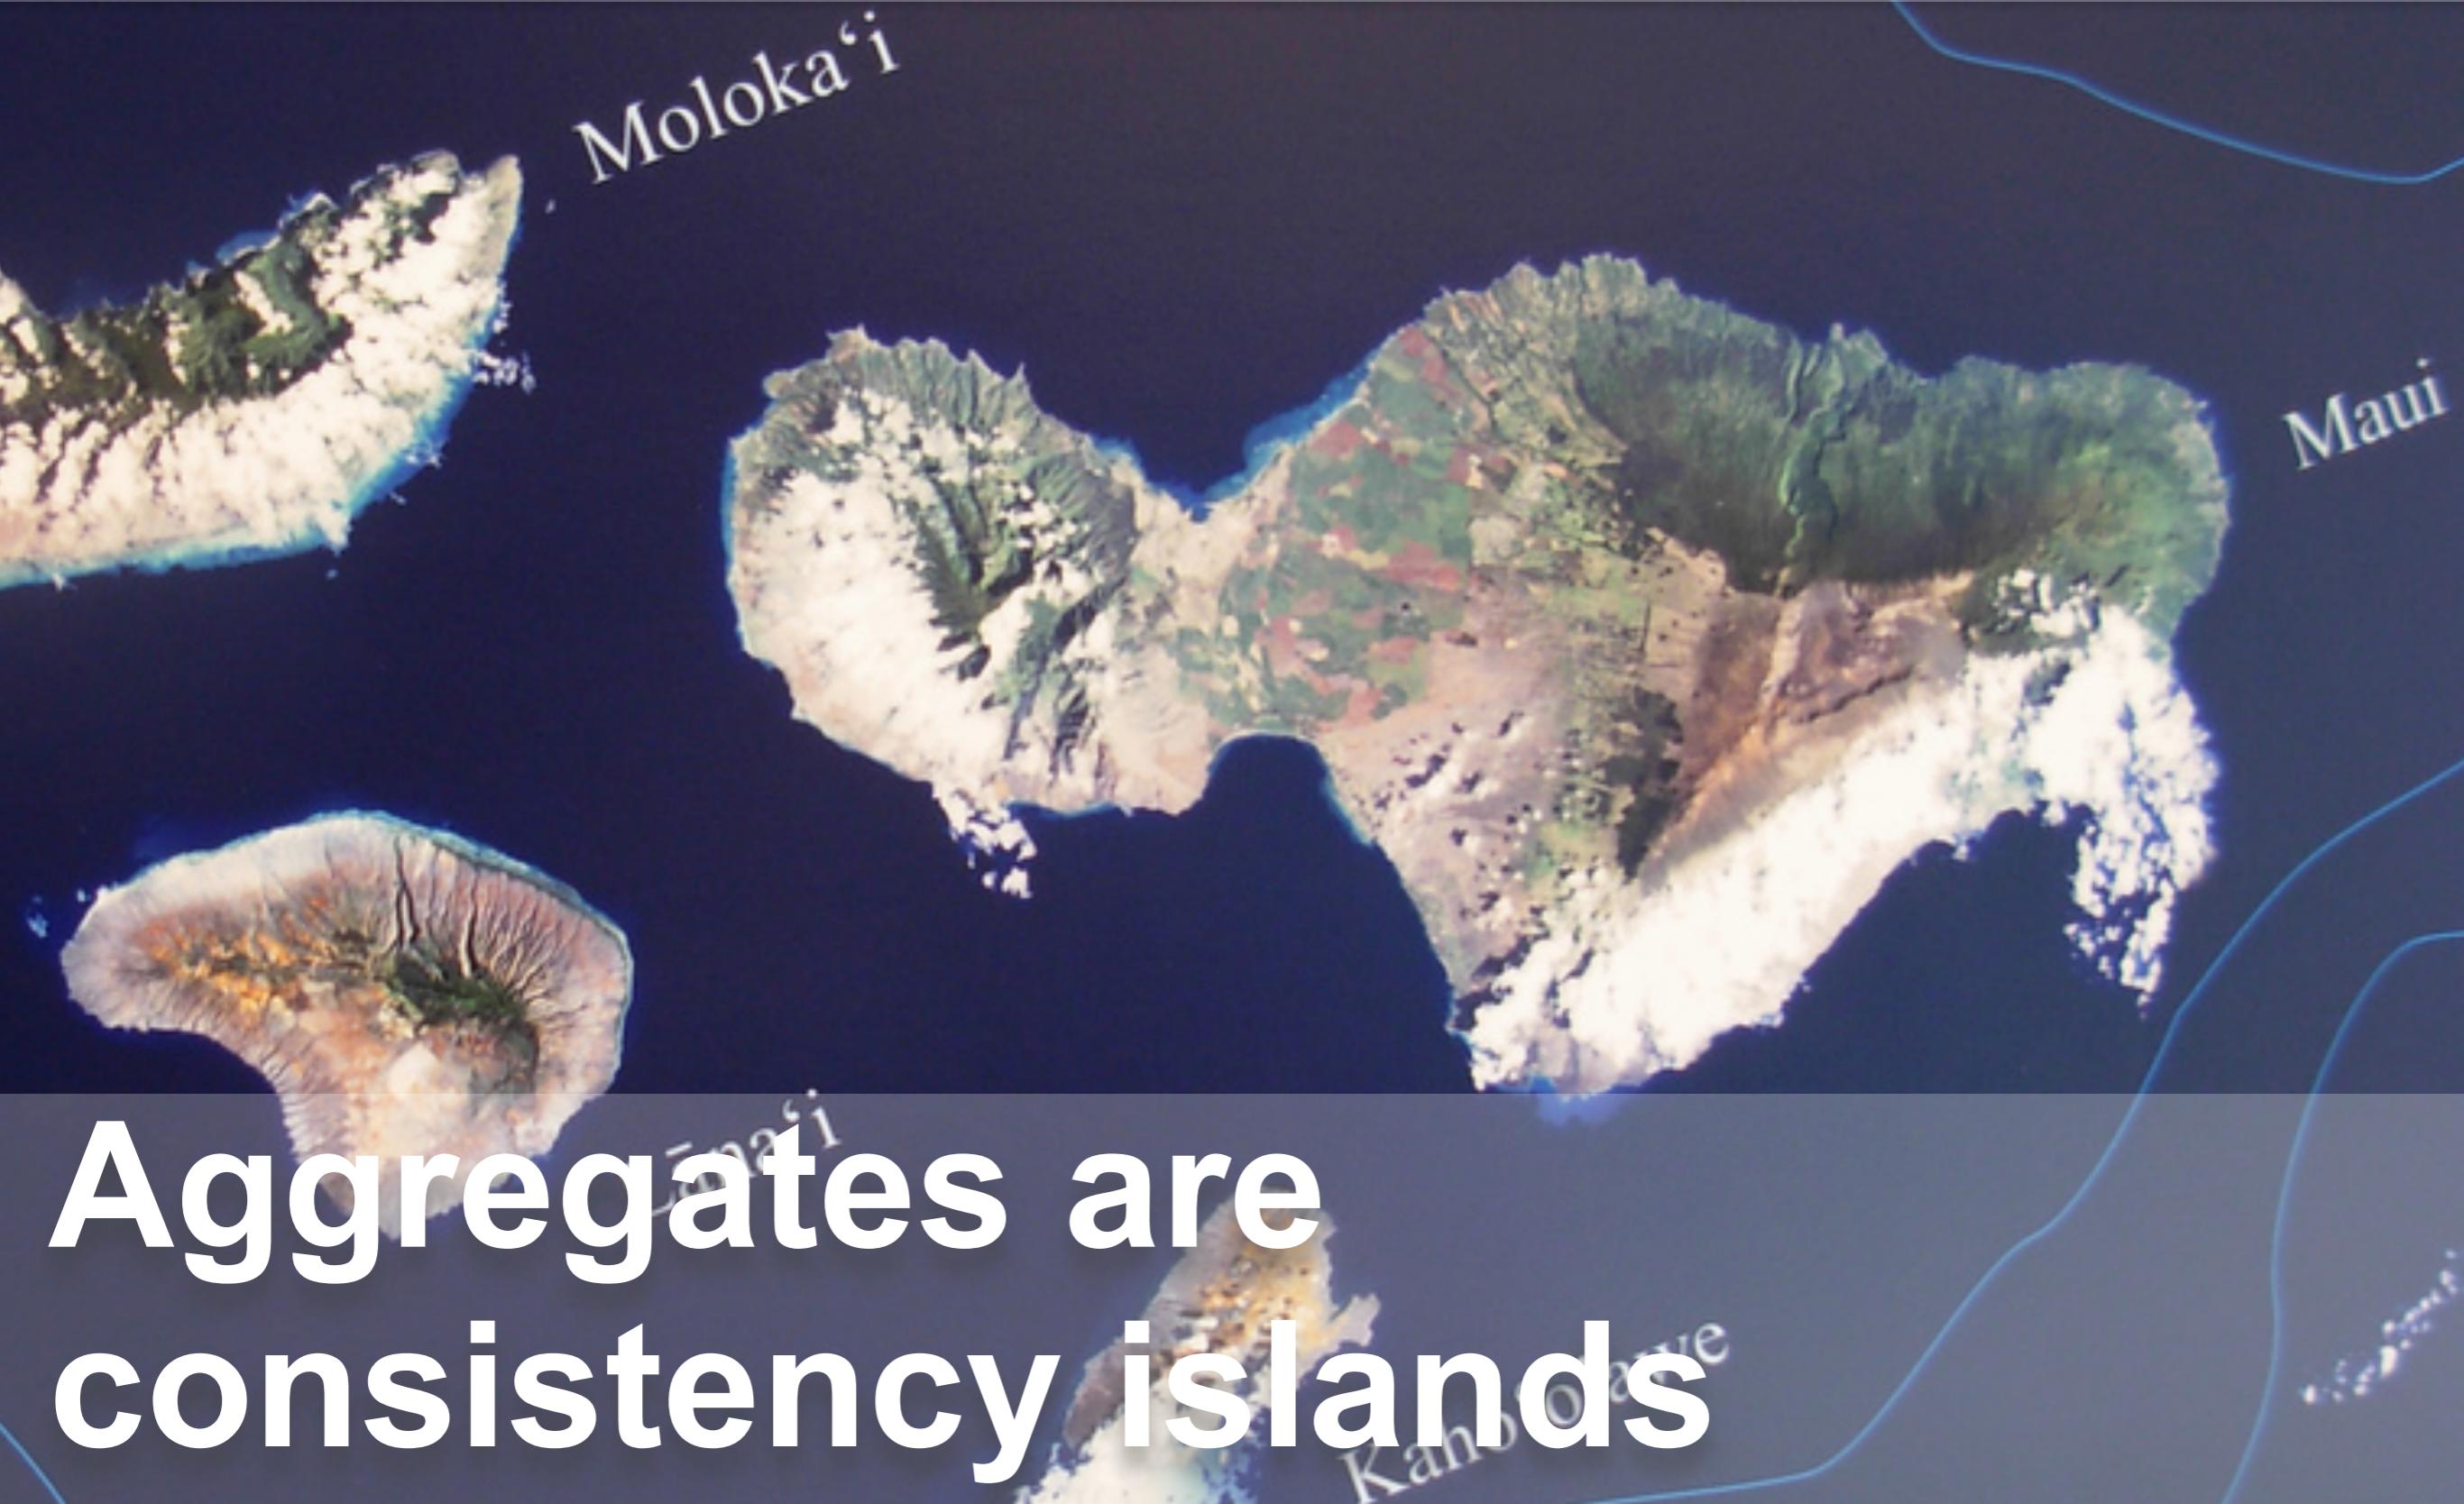
Use eventual
consistency across
aggregate boundaries

Aggregates are “eventually consistent” with each other.

- asynchronous updates propagate through system

Aggregates are always internally consistent.

Invariants apply at every transaction commit



Aggregates are
consistency islands

EventStorming

EventStorm how you think the Limit of Liability scenario should work for Repair warranties. Assume Contracts, Claims and Repair POs are different aggregates.

Include any domain events necessary, including

- Contract is terminated
- Associated open Claims are closed
- Repair POs are finalized
- Customer is notified
- Finance is notified of relevant financials



Aggregate
Persistance

Aggregates are the basic element of transfer of data storage

- you request to load or save *whole aggregates*.

Repositories are services
that provide access to
stored aggregates in terms
of the ubiquitous language

*Encapsulate actual storage
and query technology
inside repositories*

Pramodkumar J Sadalage & Martin Fowler

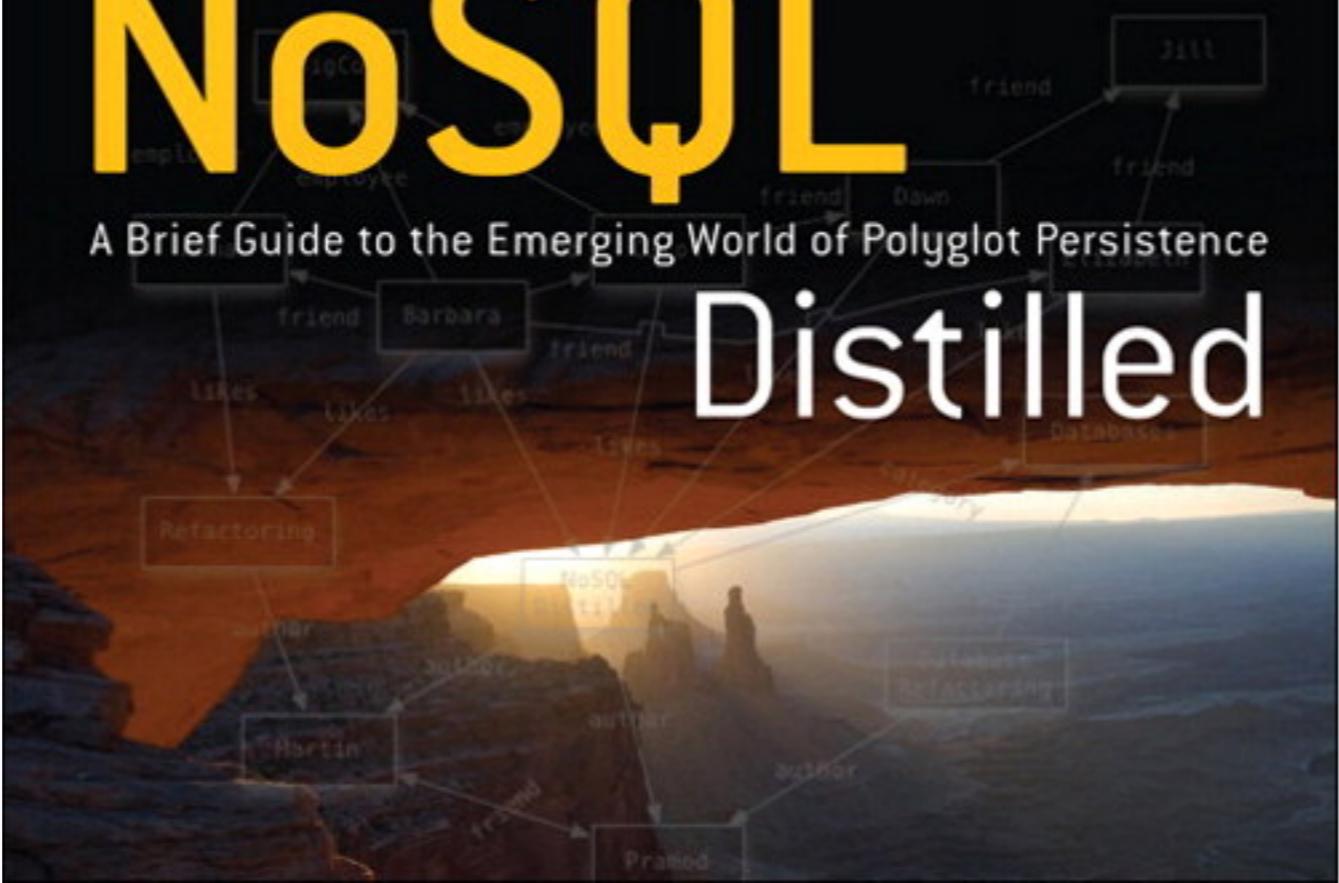
```
( "firstname": "Pramod",
  "citiesvisited": [ "Chicago", "London", "Pune", "Bangalore" ],
  "addresses": [
    { "state": "AK",
      "city": "DILLINGHAM",
      "type": "R"
    },
    { "state": "MH",
      "city": "PUNE",
      "type": "R"
    }
  ],
  "lastcity": "Chicago"
)
```

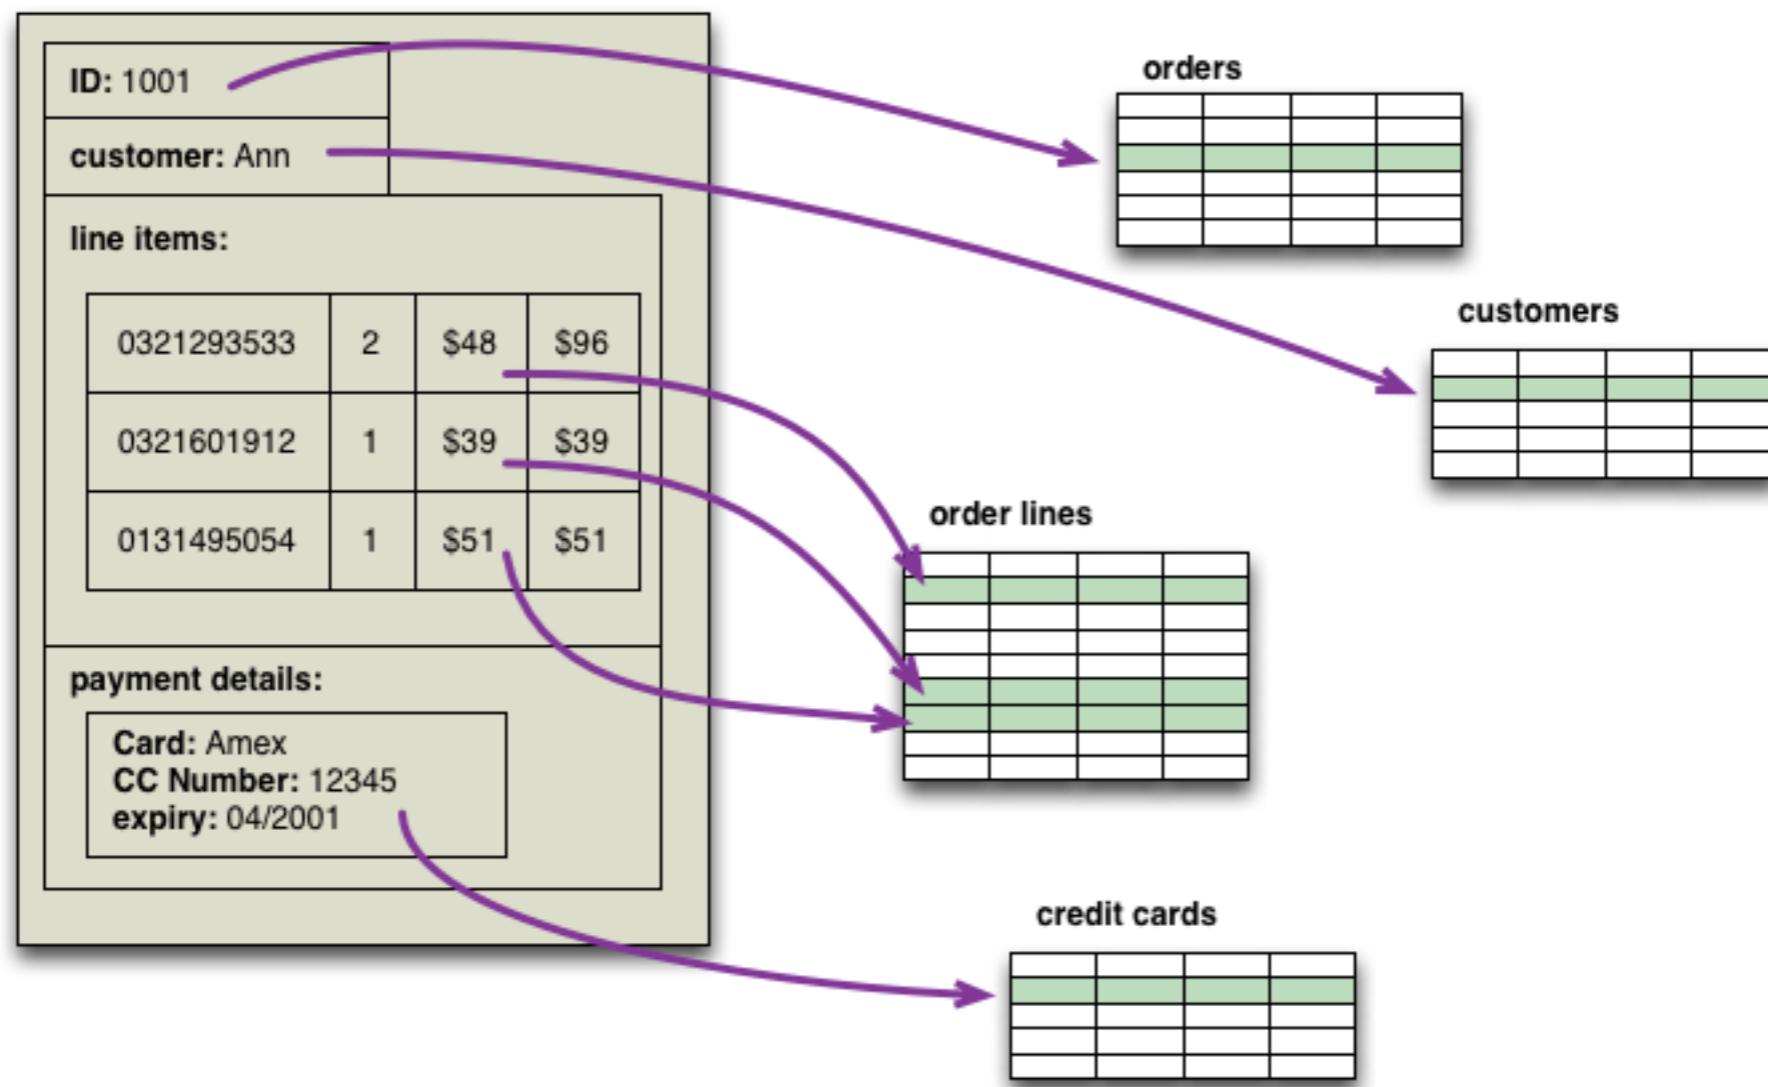


NoSQL

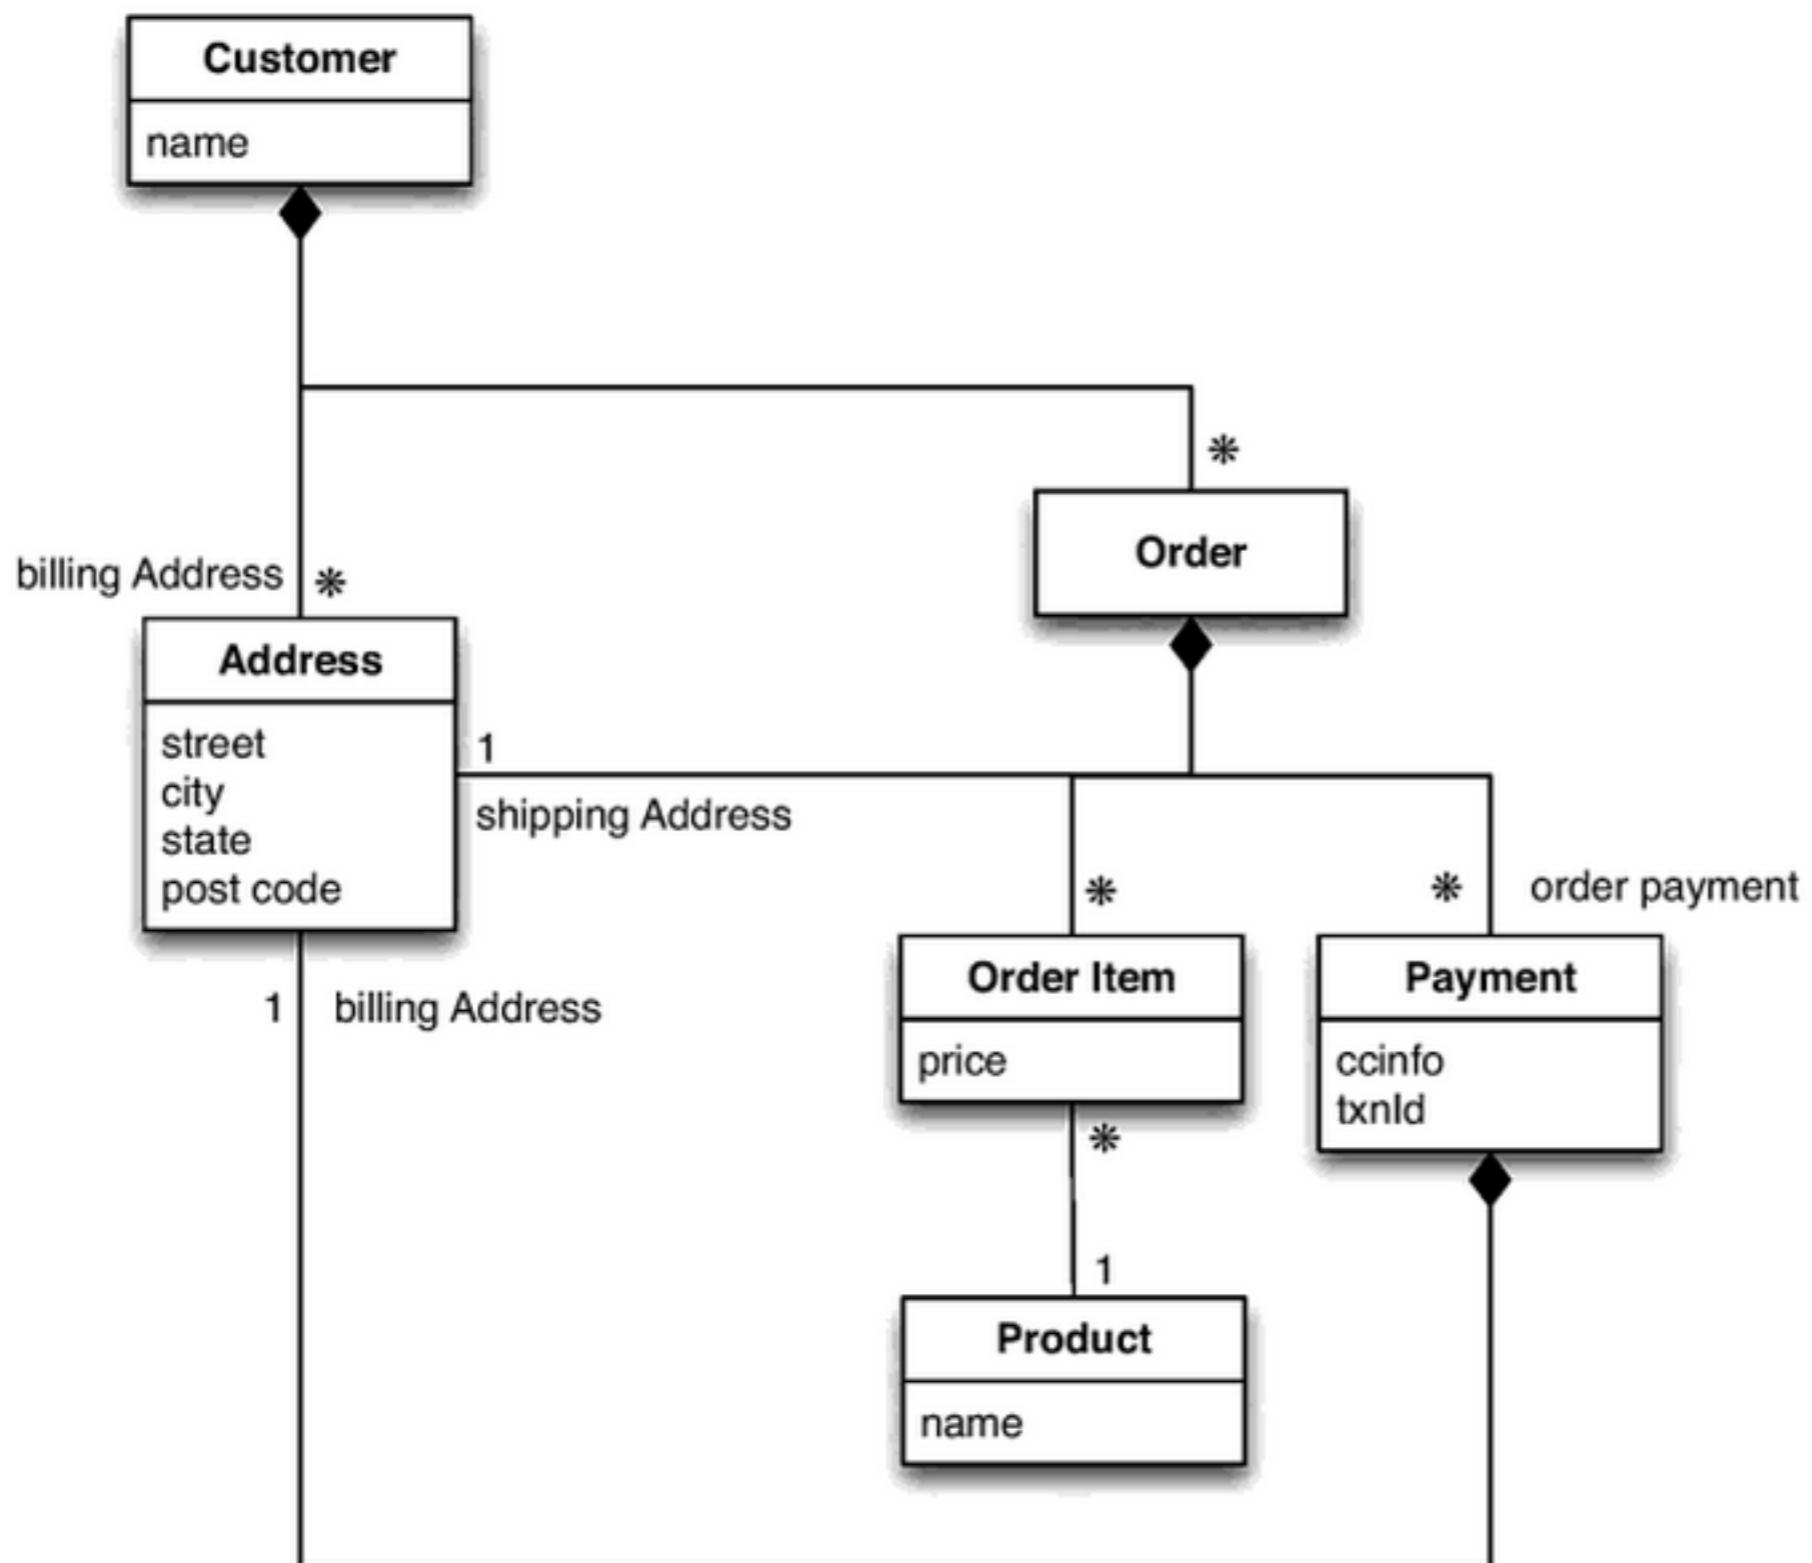
A Brief Guide to the Emerging World of Polyglot Persistence

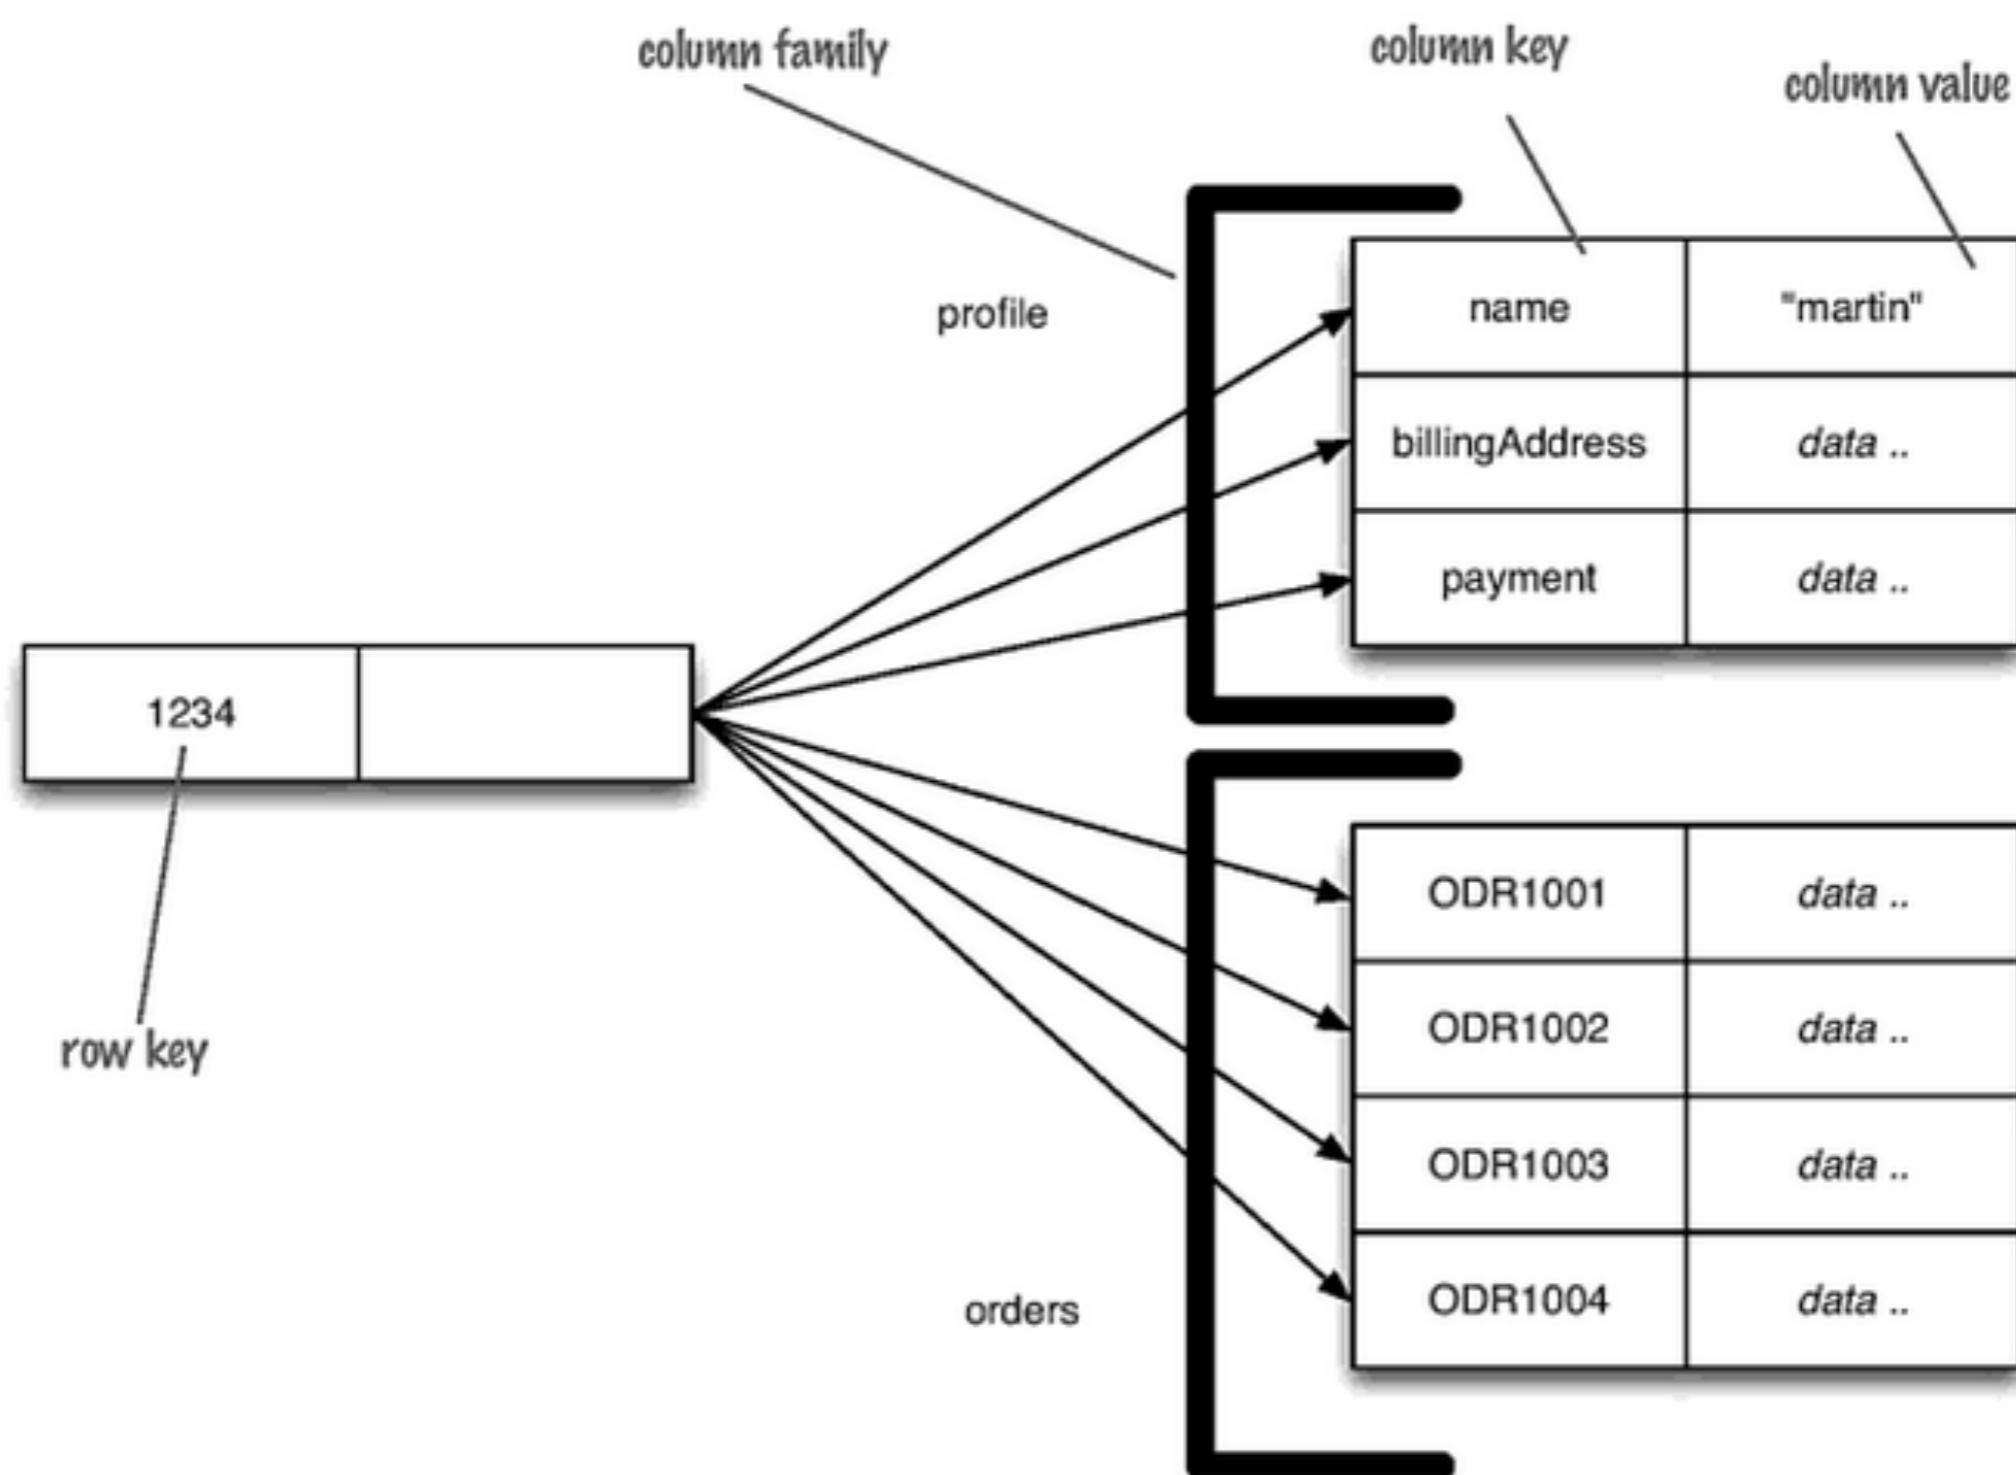
Distilled





```
{  
    orderid: '1001',  
    customerid: '232322',  
    customername: 'Ann Rogers',  
    shippingaddress: {  
        street: '1343 Broadway Street',  
        city: 'Denver',  
        zipcode: '80210',  
        state: 'CO'  
    }  
    lineitems: [  
        {id: '032193533', count: '2', cost: '$48', amount: '$96'},  
        {id: '0321601912', count: '1', cost: '$39', amount: '$39'},  
        {id: '013495054', count: '1', cost: '$51', amount: '$51'}  
    ]  
    paymentdetails : {  
        card: 'Amex',  
        ccnumber: '12345',  
        expiry: '04/2001'  
    }  
}
```





An aggregate is a collection of data that we interact with as a unit.

Aggregates form the boundaries for ACID operations with the database.



Key-value, document, and column-family databases can all be seen as forms of *aggregate-oriented database*.



Aggregates make it easier
for the database to
manage data storage over
clusters.

Aggregate-oriented databases work best

*when most data interaction is done with
the same aggregate;*

aggregate-ignorant databases are better

*when interactions use data organized in
many different formations.*